

# Infrastructure

2018 RECAP




# 2018 Big Things

**Reliability**

**Chris!**

**Ubuntu 16**

# Feature Highlights



# Feature Highlights

**Mux**

**Auth DNS**

**Zone Transfer**

- Standardized Reason Codes
- Externalized Providers
- Benefit Report
- Removed Hughes from geo rollup

YYYYMMDD

2018-10-19 ▼

▼  openmix\_customer\_hour YYYYMMDD (395)

s0. <b>queryType</b>	STRING	
s0. <b>rttScores</b>	RECORD	REPEATED
s0.rttScores. <b>reasonCode</b>	STRING	NULLABLE
s0.rttScores. <b>decisionCount</b>	INTEGER	NULLABLE
s0.rttScores. <b>scores</b>	RECORD	REPEATED
s0.rttScores.scores. <b>pid</b>	INTEGER	NULLABLE
s0.rttScores.scores. <b>value</b>	INTEGER	NULLABLE
s0. <b>tputScores</b>	RECORD	REPEATED
s0.tputScores. <b>reasonCode</b>	STRING	NULLABLE

1 Year History

Every Decision

Detail Score  
Values

Benefit report requires detailed scores for each decision.

# Changes in Apps

## Flow App & Static Failover

## Network-only Scores

- Flow app can return multiple responses
- Ignored provider API for benefit report
- Can move providers from JS code to Portal

```
1  {
2    "app_type": "V1_FLOW",
3    "app_version": 1,
4    "ttl": 20,
5    "default_cname": "no.available.candidates.g",
6    "app_flow_data": [
7      {
8        "condition": {
9          "ip": [
10             "123.45.26.23/24",
11             "2.3.4.5"
12           ]
13        },
14        "routing": [
15          {
16            "method": "FAILOVER",
17            "platforms": {
18              "platform1": {
19                "priority": 1,
20                "radar_availability": 80,
21                "response": "www.mydomain1.com",
22                "sonar": true
23              },
```


```
{
  "condition": {
    "asn": [
      30512,
      30555
    ]
  },
```

```
{
  "condition": {
    "country": [
      70,
      223
    ],
    "market": [
      3
    ]
  },
```

# Improvements to Fusion Platform

**Streaming**

**onFusion()  
callback**

- Two second fusion updates to appserver
- App can cache parsed fusion data
- No more fusion in benson 
- RMQ still SPOF





**Antoine** 🏠 11:25 AM

I did some tests and I found that on average it's taking around 1s.

---

# Features in Progress



# In Progress

**GSLB  
Rates**

**Long  
Queries**

**DNSSEC**


- DNSSEC is huge spec, lots of moving parts here. Q1 is tight 🧑🔧
- GSLB is stalled on Citrix/IT or design
- Query performance still in research 🧠



joel 🐙 1:13 PM

As the RFC puts it:

An alternating sequence of DNS public key (DNSKEY) RRsets and Delegation Signer (DS) RRsets forms a chain of signed data, with each link in the chain vouching for the next. A DNSKEY RR is used to verify the signature covering a DS RR and allows the DS RR to be authenticated. The DS RR contains a hash of another DNSKEY RR and this new DNSKEY RR is authenticated by matching the hash in the DS RR. This new DNSKEY RR in turn authenticates another DNSKEY RRset and, in turn, some DNSKEY RR in this set may be used to authenticate another DS RR, and so forth until the chain finally ends with a DNSKEY RR whose corresponding private key signs the desired DNS data.



# Backend Ready

## Canary Apps

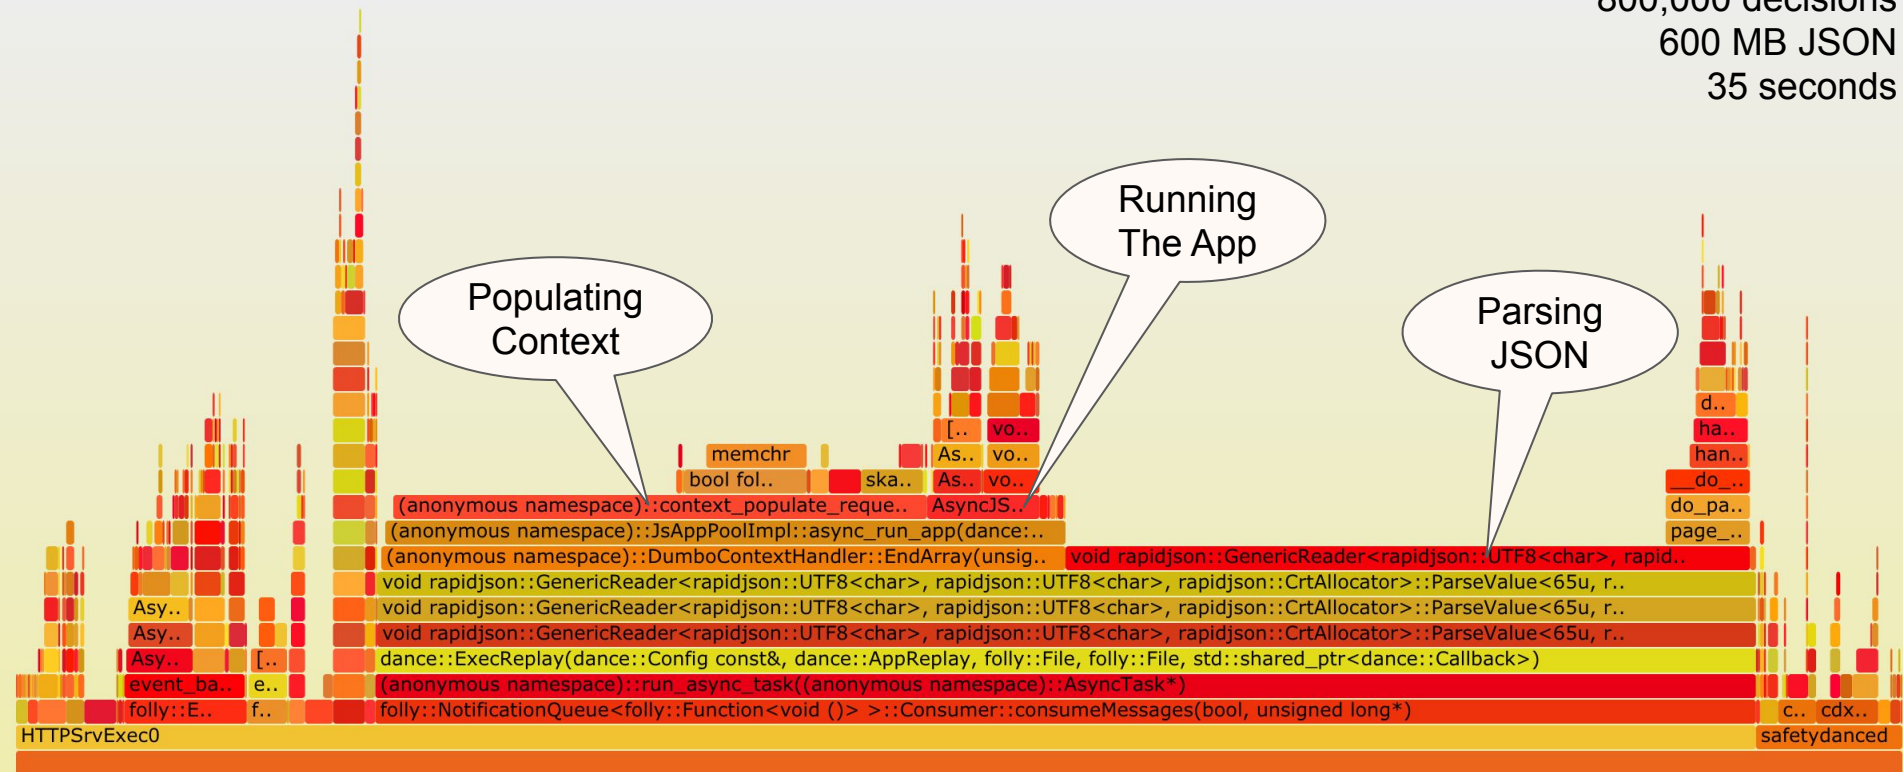
TCP  
DNS

App  
Replay

- All these are ready and pending time from UI or Ops to get it up and running.
- Sure there are bugs to be uncovered!

# Flame Graph

800,000 decisions  
600 MB JSON  
35 seconds



This is a flame graph of replay, which shows running a JS App takes approximately the same amount of time as just parsing raw JSON from Dumbo

# Platform Highlights



# Better Connections

**Autocon**

**Rate monitors**


**Reconnects**

- Moved VPN responsibility to ops
- New library code for stream management
- Per-sender stats at receiver



# Avg of appserver.score\_messages\_accepted over prod by ...


1h The Past Hour



How many scores received, as reported by each appserver

nz1.dca.sl.prod % scores dropped per ...

1h The Past Hour



Filter series

Zinter to appserver: difference of sender & receiver counts

# Reduced Administration

**Zero Config**

**No-Downtime  
Restarts**

**We run Auth DNS**

- Our own auth dns code powers [cedexis.net](https://cedexis.net)
- All upgrades are hand-over hand and downtime-free

```
"ns1.cedexis.net": { "A": { "response" : "204.48.34.10", "ttl": 86400 } },
"ns2.cedexis.net": { "A": { "response" : "209.107.211.14", "ttl": 86400 } },
"ns3.cedexis.net": { "A": { "response" : "69.28.180.4", "ttl": 86400 } },
"ns4.cedexis.net": { "A": { "response" : "172.85.88.4", "ttl": 86400 } },

"cedexis.net" : {
  "NS" : {
    "response" : [
      "flipm.cdxcn.net", "flipd.cdxcn.net", "flipg.cdxcn.net",
      "flipp.cdxcn.net"
    ],
    "ttl" : 28800
  },
  "SOA" : {
    "response": {
      "mname": "flipm.cdxcn.net",
      "rname": "admin.cedexis.com",
      "refresh": 14400,
      "retry": 7200,
      "expire": 604800,
      "rxttl": 20
    },
    "ttl": 214800
  }
}
```

---

Same format emitted by portal for customer-configured domains

## Bug Fixes & Miscellany


**Ready for  
EDNS Flag Day**

**Throughput  
scores > 64Mbps**

- JS App exceptions are now available in BQ
- Real custom query support is now in dumbo-api
- dumbo-prod-fusion is a thing that exists 🙋

# Group Varint Encoding

- Idea: encode groups of 4 values in 5-17 bytes
  - Pull out 4 2-bit binary lengths into single byte prefix



- Decode: Load prefix byte and use value to lookup in 256-entry table
  - 00000110 → ...  
Offsets: +1,+2,+3,+5; Masks: ff, ff, ffff  
...
- Much faster than alternatives:
  - 7-bit-per-byte varint: decode ~180M numbers/second
  - 30-bit Varint w/ 2-bit length: decode ~240M numbers/second
  - Group varint: decode ~400M numbers/second

Row	byte_size	pct
1	1 byte	2.15
2	2 byte	78.18
3	3 byte	19.67

- Was uint16\_t which was max at 64Mbps
- Most values still take 2 bytes

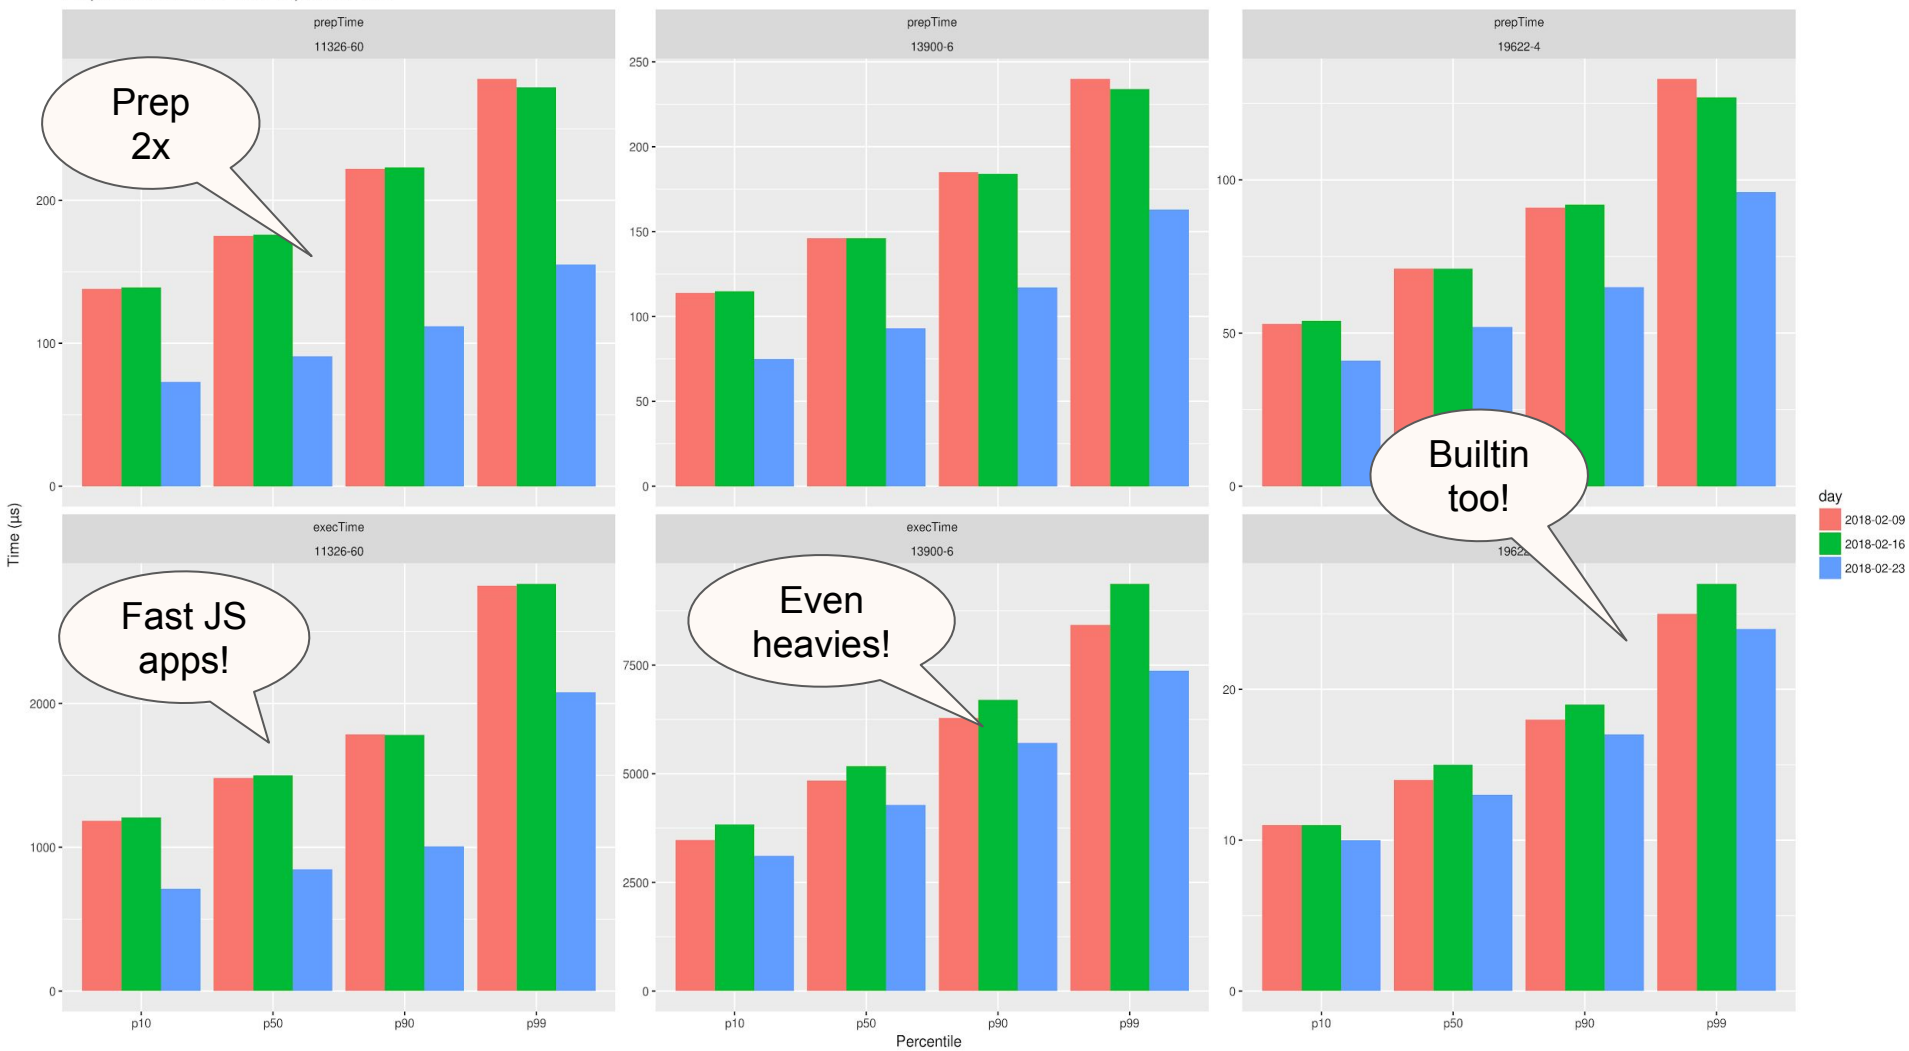
# Appserver Performance Improvements

## Immutable Apps

## JSProc IPC Serialization

- Zero-allocation trivially copyable structs
- Removed runtime checks in apps; they now reload when config changes

# Prep Time and Exec Time improvements





## Other Performance Improvements

**150k/req/sec**  
**DNS**

**3x URL parsing**  
**for reports**

- Netmap parse shaved 20sec from boot
- New stats lib with no cross-thread sharing
- New rate limit implementations
- Faster logging too!

```

297 // N2
298 template <>
299 void parse_path<n1_report>(RadarLog& log, folly::StringPiece path) {
300     using NavTiming = openmix::messages::RadarNavTimingValues;
301     log.set_report_type(ReportType::N1Report);
302
303     PathParse p = {path};
304     start_parse(p, log, "/n1")           //
305         / &RadarLog::set_response_code //
306         ;
307     parse_into(p, *log.mutable_navtiming())
308         / &NavTiming::set_navigation_start //
309         / &NavTiming::set_unload_event_start //
310         / &NavTiming::set_unload_event_end //
311         / &NavTiming::set_redirect_start
312         / &NavTiming::set_redirect_end
313         / &NavTiming::set_fetch_start
314         / &NavTiming::set_domain_lookup_start
315         ;
316     parse_into(p, log)
317         / required(&RadarLog::mutable_signature)
318         / optional(&RadarLog::mutable_partner_tags)
319         / optional(&RadarLog::set_start_render)
320         ;
321     validate_parse(log, p.ok);
322 }

```

1

2

```

template <typename T, typename FieldT>
ParseContext<T>& operator/(ParseContext<T>& ctx, FieldT&& fld) {
    parse_segment(ctx, fld, ctx.parse.path.split_step('/'));
    return ctx;
}

```

3

```

#define CALL_MEMBER_FN(object, ptrToMember) ((object).*(ptrToMember))
template <typename T, typename V>
void parse_segment(ParseContext<T>& ctx, IntegralField<T, V> fld,
                  folly::StringPiece segment) {
    if (ctx.parse.ok) {
        if (auto v = folly::tryTo<V>(segment)) {
            CALL_MEMBER_FN(ctx.target, fld.setter)(ctx.basis + *v);
        } else if (fld.required) {
            ctx.parse.ok = false;
        }
    }
}

```

# Dumbo Performance Improvements

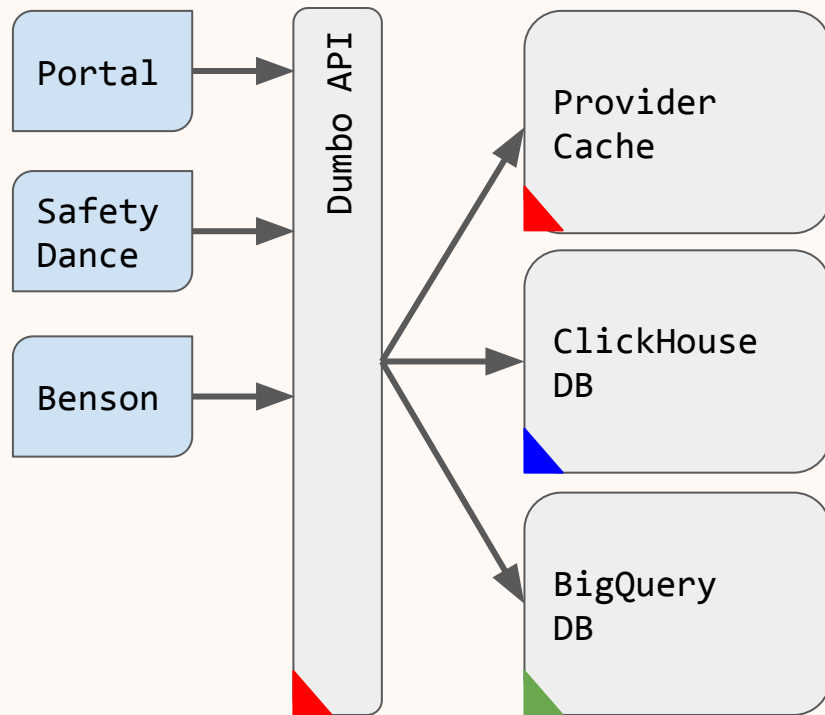

**Gzip responses**

**Last-hour API  
queries are  
sub-second**

**30 day BQ works!**

- New CH hardware, better SQL
- Better BQ SQL
- Still super reliable
- Ingest performance is ~100s

# Dumbo API Architecture Review



# Team Highlights



# Disciplined Release Schedule

**Frequent  
Releases**

**184 Releases  
163 Hosts**

appserver: 26

autocon: 3

belvedere: 3

benson: 10

button: 14

deliverator: 1

dumbo-api: 24

dumbo-ingest: 9

dumbo-notify: 4

kennylogger: 21

newzinter: 12

safety-dance: 16

spinderella: 5

tangled: 7

tps-report: 11

tribble: 2

# releases - Jan 2nd



**darren** 🇬🇧 11:28 AM

wow

# releases - Dec 14th, 2018



**alok** 10:27 AM

wow

# releases - Mar 12th, 2018



**darren** 🇬🇧 10:34 AM

wow, ok!

# releases - Feb 22nd, 2018



**alok** 12:39 PM

wow, huge release. 👍

# releases - Sep 13th, 2018



**alok** 10:07 AM

YES! HUGE release!



# Code Highlights

**DNS Server  
Rewrite**

**New Unit Test  
Library**

- Stats lib
- Builtins lib
- Send/Recv lib
- CSV  $\Rightarrow$  JSON config
- Direct apt deps



# Changes in Openmix2.0 in 2018



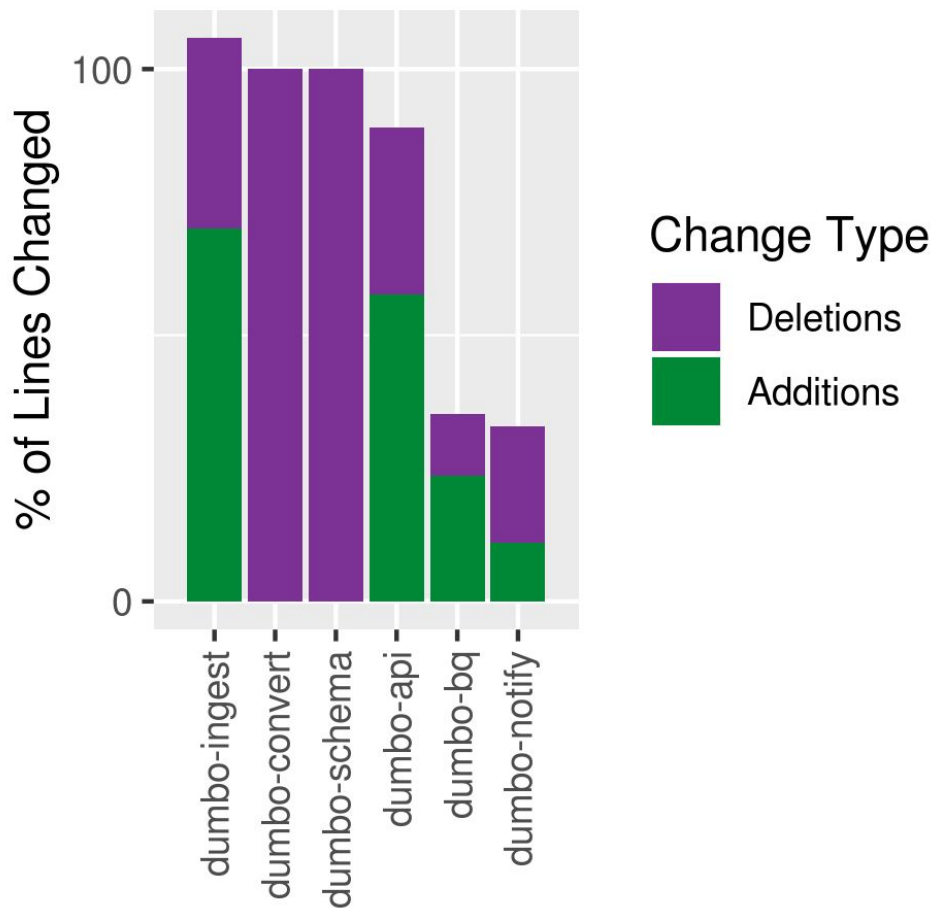
This code churn chart is the best chart in the entire world!

# Healthy Churn

Safe releases  
Safe changes  
Safe culture

Everybody is at  
the Party

## Changes in Bigdata in 2018



# Team Safety

















## How We Do It

**No Bug  
Backlog**

**No JIRA**

- Big Active Project
- Individuals keep lists of things which bug them
- Trusted to fix those things
- 出る釘は打たれる

<input type="checkbox"/>	 <b>dnsserver: Add code to generate RRSIGs</b> #2196 by cincodenada was merged on Dec 12, 2018 • Approved	 86
<input type="checkbox"/>	 <b>dnsserver: Rework ECS/EDNS handling</b> #2134 by cincodenada was merged on Nov 9, 2018 • Approved	 63
<input type="checkbox"/>	 <b>DNSServer: listen TCP</b> #2057 by suspend0 was merged on Oct 11, 2018 • Approved	 94
<input type="checkbox"/>	 <b>Safety: app replay feature [WIP]</b> #2019 by suspend0 was merged on Sep 25, 2018 • Approved	 131
<input type="checkbox"/>	 <b>Providers: single thread for all consumers</b> #1795 by suspend0 was merged on Jun 12, 2018 • Approved	 52
<input type="checkbox"/>	 <b>JSProc: Async implementation</b> #1687 by suspend0 was merged on May 11, 2018 • Approved	 64
<input type="checkbox"/>	 <b>Appserver: new multi-stage app</b> #1653 by suspend0 was merged on Apr 21, 2018 • Approved	 107



Building consensus about the change: answer questions & debate implementation  
Some pull requests don't get enough support and are abandoned

# Kaizen (改善) and Sekinin (責任)

**Team  
Responsibility**

**Small  
Improvements**

- Common to have 20/50/100 comments on a code review
- Safe to ask a question
- Safe to ask why
- Safe to try something
- Safe to be wrong

“If you are going to do kaizen continuously...you’ve got to assume that things are a mess. Too many people just assume that things are alright the way they are.

Are you guys convinced that the way you’re doing things is the right way? That’s no way to get anything done.

Kaizen is about changing the way things are. If you assume that things are alright the way they are, you can’t do kaizen. So change something!”

–Taiichi Ohno

# Practicing Continuous Improvement

**Have a  
Standard**

**Change the  
Standard**

- Monorepo
- Shared code
- Copy & tweak
- Tension between  
change & consistency



# Team Safety

**Messy is OK**

**Temporary is OK**

**Fail is OK**

- Get something in, get it shipped
- Make better later
- But, actually do that

fin